

```
In [1]: import gzip
import math
import random
from collections import defaultdict
```

Data is available at <http://cseweb.ucsd.edu/~jmcauley/pml/data/>. Download and save to your own directory

```
In [2]: dataDir = "/home/jmcauley/pml_data/"
```

Amazon musical instrument review data. Originally from <https://s3.amazonaws.com/amazon-reviews-pds/tsv/index.txt>

```
In [3]: path = dataDir + "amazon_reviews_us_Musical_Instruments_v1_00.tsv.gz"
f = gzip.open(path, 'rt', encoding="utf8")

header = f.readline()
header = header.strip().split('\t')
```

Dataset contains the following fields

```
In [4]: header
```

```
Out[4]: ['marketplace',
'customer_id',
'review_id',
'product_id',
'product_parent',
'product_title',
'product_category',
'star_rating',
'helpful_votes',
'total_votes',
'vine',
'verified_purchase',
'review_headline',
'review_body',
'review_date']
```

Parse the data and convert fields to integers where needed

```
In [5]: dataset = []

for line in f:
    fields = line.strip().split('\t')
    d = dict(zip(header, fields))
    d['star_rating'] = int(d['star_rating'])
    d['helpful_votes'] = int(d['helpful_votes'])
    d['total_votes'] = int(d['total_votes'])
    dataset.append(d)
```

One row of the dataset (as a python dictionary)

In [6]: `dataset[0]`

```
Out[6]: {'customer_id': '45610553',
        'helpful_votes': 0,
        'marketplace': 'US',
        'product_category': 'Musical Instruments',
        'product_id': 'B00HH62VB6',
        'product_parent': '618218723',
        'product_title': 'AGPtek® 10 Isolated Output 9V 12V 18V Guitar Pedal Board Power Supply Effect Pedals with Isolated Short Circuit / Overcurrent Protection',
        'review_body': 'Works very good, but induces ALOT of noise.',
        'review_date': '2015-08-31',
        'review_headline': 'Three Stars',
        'review_id': 'RMDCHWD0Y5OZ9',
        'star_rating': 3,
        'total_votes': 1,
        'verified_purchase': 'N',
        'vine': 'N'}
```

Extract a few utility data structures

```
In [7]: usersPerItem = defaultdict(set) # Maps an item to the users who rated it
        itemsPerUser = defaultdict(set) # Maps a user to the items that they rated
        itemNames = {}
        ratingDict = {} # To retrieve a rating for a specific user/item pair

        for d in dataset:
            user,item = d['customer_id'], d['product_id']
            usersPerItem[item].add(user)
            itemsPerUser[user].add(item)
            ratingDict[(user,item)] = d['star_rating']
            itemNames[item] = d['product_title']
```

Extract per-user and per-item averages (useful later for rating prediction)

```
In [8]: userAverages = {}
        itemAverages = {}

        for u in itemsPerUser:
            rs = [ratingDict[(u,i)] for i in itemsPerUser[u]]
            userAverages[u] = sum(rs) / len(rs)

        for i in usersPerItem:
            rs = [ratingDict[(u,i)] for u in usersPerItem[i]]
            itemAverages[i] = sum(rs) / len(rs)
```

## Similarity metrics

### Jaccard

```
In [9]: def Jaccard(s1, s2):
        numer = len(s1.intersection(s2))
        denom = len(s1.union(s2))
        if denom == 0:
```

```

    return 0
    return numer / denom

```

## Cosine

Simple implementation for set-structured data

```

In [10]: def CosineSet(s1, s2):
    # Not a proper implementation, operates on sets so correct for interactions
    numer = len(s1.intersection(s2))
    denom = math.sqrt(len(s1)) * math.sqrt(len(s2))
    if denom == 0:
        return 0
    return numer / denom

```

Or for real values (e.g. ratings). Note that this implementation uses global variables (usersPerItem, ratingDict), which ideally should be passed as parameters.

```

In [11]: def Cosine(i1, i2):
    # Between two items
    inter = usersPerItem[i1].intersection(usersPerItem[i2])
    numer = 0
    denom1 = 0
    denom2 = 0
    for u in inter:
        numer += ratingDict[(u,i1)]*ratingDict[(u,i2)]
    for u in usersPerItem[i1]:
        denom1 += ratingDict[(u,i1)]**2
    for u in usersPerItem[i2]:
        denom2 += ratingDict[(u,i2)]**2
    denom = math.sqrt(denom1) * math.sqrt(denom2)
    if denom == 0: return 0
    return numer / denom

```

## Pearson

```

In [12]: def Pearson(i1, i2):
    # Between two items
    iBar1 = itemAverages[i1]
    iBar2 = itemAverages[i2]
    inter = usersPerItem[i1].intersection(usersPerItem[i2])
    numer = 0
    denom1 = 0
    denom2 = 0
    for u in inter:
        numer += (ratingDict[(u,i1)] - iBar1)*(ratingDict[(u,i2)] - iBar2)
    for u in inter: #usersPerItem[i1]:
        denom1 += (ratingDict[(u,i1)] - iBar1)**2
    #for u in usersPerItem[i2]:
        denom2 += (ratingDict[(u,i2)] - iBar2)**2
    denom = math.sqrt(denom1) * math.sqrt(denom2)
    if denom == 0: return 0
    return numer / denom

```

# Retrieve the most similar items to a given query

In this case, based on the Jaccard similarity

```
In [13]: def mostSimilar(i, N):
          similarities = []
          users = usersPerItem[i]
          for i2 in usersPerItem:
              if i2 == i: continue
              sim = Jaccard(users, usersPerItem[i2])
              #sim = Pearson(i, i2) # Could use alternate similarity metrics straightf
              similarities.append((sim,i2))
          similarities.sort(reverse=True)
          return similarities[:10]
```

Choose an item to use as a query

```
In [14]: dataset[2]
```

```
Out[14]: {'customer_id': '6111003',
          'helpful_votes': 0,
          'marketplace': 'US',
          'product_category': 'Musical Instruments',
          'product_id': 'B0006VMBHI',
          'product_parent': '603261968',
          'product_title': 'AudioQuest LP record clean brush',
          'review_body': 'removes dust. does not clean',
          'review_date': '2015-08-31',
          'review_headline': 'Three Stars',
          'review_id': 'RIZR67JKUDBI0',
          'star_rating': 3,
          'total_votes': 1,
          'verified_purchase': 'Y',
          'vine': 'N'}
```

```
In [15]: query = dataset[2]['product_id']
```

Retrieve the most similar items

```
In [16]: ms = mostSimilar(query, 10)
```

```
In [17]: ms
```

```
Out[17]: [(0.028446389496717725, 'B00006I5SD'),
          (0.01694915254237288, 'B00006I5SB'),
          (0.015065913370998116, 'B000AJR482'),
          (0.014204545454545454, 'B00E7MVP3S'),
          (0.008955223880597015, 'B001255YL2'),
          (0.008849557522123894, 'B003EIRVO8'),
          (0.008333333333333333, 'B0015VEZ22'),
          (0.00821917808219178, 'B00006I5UH'),
          (0.008021390374331552, 'B00008BWM7'),
          (0.007656967840735069, 'B000H2BC4E')]
```

Print names of query and recommended items

```
In [18]: itemNames[query]
```

```
Out[18]: 'AudioQuest LP record clean brush'
```

```
In [19]: [itemNames[x[1]] for x in ms]
```

```
Out[19]: ['Shure SFG-2 Stylus Tracking Force Gauge',
'Shure M97xE High-Performance Magnetic Phono Cartridge',
'ART Pro Audio DJPRE II Phono Turntable Preamplifier',
'Signstek Blue LCD Backlight Digital Long-Playing LP Turntable Stylus Force Scale Gauge Tester',
'Audio Technica AT120E/T Standard Mount Phono Cartridge',
'Technics: 45 Adaptor for Technics 1200 (SFWE010)',
'GruvGlide GRUVGLIDE DJ Package',
'STANTON MAGNETICS Record Cleaner Kit',
'Shure M97xE High-Performance Magnetic Phono Cartridge',
'Behringer PP400 Ultra Compact Phono Preamplifier']
```

Faster implementation

```
In [20]: def mostSimilarFast(i, N):
similarity = []
users = usersPerItem[i]
candidateItems = set()
for u in users:
    candidateItems = candidateItems.union(itemsPerUser[u])
for i2 in candidateItems:
    if i2 == i: continue
    sim = Jaccard(users, usersPerItem[i2])
    similarity.append((sim,i2))
similarity.sort(reverse=True)
return similarity[:N]
```

Confirm that results are the same...

```
In [21]: mostSimilarFast(query, 10)
```

```
Out[21]: [(0.028446389496717725, 'B00006I5SD'),
(0.01694915254237288, 'B00006I5SB'),
(0.015065913370998116, 'B000AJR482'),
(0.014204545454545454, 'B00E7MVP3S'),
(0.008955223880597015, 'B001255YL2'),
(0.008849557522123894, 'B003EIRVO8'),
(0.008333333333333333, 'B0015VEZ22'),
(0.00821917808219178, 'B00006I5UH'),
(0.008021390374331552, 'B00008BWM7'),
(0.007656967840735069, 'B000H2BC4E')]
```

## Similarity-based rating estimation

Use our similarity functions to estimate ratings. Start by building a few utility data structures.

```
In [22]: reviewsPerUser = defaultdict(list)
reviewsPerItem = defaultdict(list)
```

```
In [23]: for d in dataset:
user,item = d['customer_id'], d['product_id']
reviewsPerUser[user].append(d)
reviewsPerItem[item].append(d)
```

```
In [24]: ratingMean = sum([d['star_rating'] for d in dataset]) / len(dataset)
```

```
In [25]: ratingMean
```

```
Out[25]: 4.251102772543146
```

Rating prediction heuristic (several alternatives from Chapter 4 could be used)

```
In [26]: def predictRating(user,item):
ratings = []
similarities = []
for d in reviewsPerUser[user]:
i2 = d['product_id']
if i2 == item: continue
ratings.append(d['star_rating'] - itemAverages[i2])
similarities.append(Jaccard(usersPerItem[item],usersPerItem[i2]))
if (sum(similarities) > 0):
weightedRatings = [(x*y) for x,y in zip(ratings,similarities)]
return itemAverages[item] + sum(weightedRatings) / sum(similarities)
else:
# User hasn't rated any similar items
return ratingMean
```

```
In [27]: dataset[1]
```

```
Out[27]: {'customer_id': '14640079',
'helpful_votes': 0,
'marketplace': 'US',
'product_category': 'Musical Instruments',
'product_id': 'B003LRN53I',
'product_parent': '986692292',
'product_title': 'Sennheiser HD203 Closed-Back DJ Headphones',
'review_body': 'Nice headphones at a reasonable price.',
'review_date': '2015-08-31',
'review_headline': 'Five Stars',
'review_id': 'RZSL0BALIYUNU',
'star_rating': 5,
'total_votes': 0,
'verified_purchase': 'Y',
'vine': 'N'}
```

Predict a rating for a particular user/item pair

```
In [28]: u,i = dataset[1]['customer_id'], dataset[1]['product_id']
```

```
In [29]: predictRating(u, i)
```

```
Out[29]: 4.509357030989021
```

Compute the MSE for a model based on this heuristic

```
In [30]: def MSE(predictions, labels):  
         differences = [(x-y)**2 for x,y in zip(predictions,labels)]  
         return sum(differences) / len(differences)
```

Compared to a trivial predictor which always predicts the mean

```
In [31]: alwaysPredictMean = [ratingMean for d in dataset]
```

Get predictions for all instances (fairly slow!)

```
In [32]: simPredictions = [predictRating(d['customer_id'], d['product_id']) for d in data]
```

```
In [33]: labels = [d['star_rating'] for d in dataset]
```

```
In [34]: MSE(alwaysPredictMean, labels)
```

```
Out[34]: 1.4796142779564334
```

```
In [35]: MSE(simPredictions, labels)
```

```
Out[35]: 1.44672577948388
```

```
In [ ]:
```

## Exercises

### 4.1

(implementation is provided via the function mostSimilarFast above)

### 4.2

(using Amazon musical instruments data from examples above)

```
In [36]: def simTest(simFunction, nUserSamples):  
         sims = []  
         randomSims = []  
  
         items = set(usersPerItem.keys())  
         users = list(itemsPerUser.keys())
```

```

for u in random.sample(users, nUserSamples):
    itemsU = set(itemsPerUser[u])
    if len(itemsU) < 2: continue # User needs at least two interactions
    (i,j) = random.sample(itemsU, 2)
    k = random.sample(items.difference(itemsU),1)[0]
    usersi = usersPerItem[i].difference(set([u]))
    usersj = usersPerItem[j].difference(set([u]))
    usersk = usersPerItem[k].difference(set([u]))
    sims.append(simFunction(usersi,usersj))
    randomSims.append(simFunction(usersi,usersk))

print("Average similarity = " + str(sum(sims)/len(sims)))
print("Average similarity (with random item) = " + str(sum(randomSims)/len(r

```

In [37]:

```
simTest(Jaccard, 1000)
```

```

Average similarity = 0.0019330961239460492
Average similarity (with random item) = 0.0

```

In [38]:

```
simTest(CosineSet, 1000)
```

```

Average similarity = 0.005634438126569325
Average similarity (with random item) = 0.0

```

## 4.3

In [39]:

```

items = set(usersPerItem.keys())
users = set(itemsPerUser.keys())

```

In [40]:

```

# 1: Average cosine similarity between i and items in u's history
def reclscore(u, i, userHistory):
    if len(userHistory) == 0:
        return 0
    averageSim = []
    s1 = usersPerItem[i].difference(set([u]))
    for h in userHistory:
        s2 = usersPerItem[h].difference(set([u]))
        averageSim.append(Jaccard(s1,s2))
    averageSim = sum(averageSim)/len(averageSim)
    return averageSim

# 2: Jaccard similarity with most similar user who has consumed i
def rec2score(u, i, userHistory):
    bestSim = None
    for v in usersPerItem[i]:
        if u == v:
            continue
        sim = Jaccard(userHistory, itemsPerUser[v])
        if bestSim == None or sim > bestSim:
            bestSim = sim
    if bestSim == None:
        return 0
    return bestSim

```



```
# Generate a recommendation for a user based on a given scoring function
def rec(u, score):
    history = itemsPerUser[u]
    if len(history) > 5: # If the history is too long, just take a sample
        history = random.sample(history,5)
    bestItem = None
    bestScore = None
    for i in items:
        if i in itemsPerUser[u]: continue
        s = score(u, i, history)
        if bestItem == None or s > bestScore:
            bestItem = i
            bestScore = s

    return bestItem, bestScore
```

```
In [41]: u = random.sample(users,1)[0]
```

```
In [42]: rec(u, rec1score)
```

```
Out[42]: ('B002KYLGT8', 0.043478260869565216)
```

```
In [43]: rec(u, rec2score)
```

```
Out[43]: ('B00HCPTXJA', 0.5)
```

```
In [44]: def recTest(simFunction, nUserSamples):
    items = set(usersPerItem.keys())
    users = list(itemsPerUser.keys())

    better = 0
    worse = 0

    for u in random.sample(users, nUserSamples):
        itemsU = set(itemsPerUser[u])
        if len(itemsU) < 2:
            continue
        i = random.sample(itemsU, 1)[0]
        uWithheld = itemsU.difference(set([i]))
        j = random.sample(items,1)[0]

        si = simFunction(u,i,uWithheld)
        sj = simFunction(u,j,uWithheld)

        if si > sj:
            better += 1
        if sj > si:
            worse += 1

    print("Better than random " + str(better) + " times")
    print("Worse than random " + str(worse) + " times")
```

Results on this dataset aren't particularly interesting. Could try with a denser dataset (so that

many items have non-zero similarity) to get more interesting results.

In [45]:

```
recTest(rec1score,5000)
```

Better than random 306 times  
Worse than random 1 times

In [46]:

```
recTest(rec2score,5000)
```

Better than random 278 times  
Worse than random 4 times

## 4.4

(following code and auxiliary data structures from the examples above)

Equation 4.20

In [47]:

```
def predictRating1(user,item):
    ratings = []
    similarities = []
    for d in reviewsPerUser[user]:
        i2 = d['product_id']
        if i2 == item: continue
        ratings.append(d['star_rating'])
        similarities.append(Jaccard(usersPerItem[item],usersPerItem[i2]))
    if (sum(similarities) > 0):
        weightedRatings = [(x*y) for x,y in zip(ratings,similarities)]
        return sum(weightedRatings) / sum(similarities)
    else:
        return ratingMean
```

Equation 4.21

In [48]:

```
def predictRating2(user,item):
    ratings = []
    similarities = []
    for d in reviewsPerItem[item]:
        u2 = d['customer_id']
        if u2 == user: continue
        ratings.append(d['star_rating'])
        similarities.append(Jaccard(itemsPerUser[user],itemsPerUser[u2]))
    if (sum(similarities) > 0):
        weightedRatings = [(x*y) for x,y in zip(ratings,similarities)]
        return sum(weightedRatings) / sum(similarities)
    else:
        return ratingMean
```

Equation 4.22

In [49]:

```
def predictRating3(user,item):
    ratings = []
    similarities = []
    for d in reviewsPerUser[user]:
        i2 = d['product_id']
```

```
    if i2 == item: continue
    ratings.append(d['star_rating'] - itemAverages[i2])
    similarities.append(Jaccard(usersPerItem[item],usersPerItem[i2]))
    if (sum(similarities) > 0):
        weightedRatings = [(x*y) for x,y in zip(ratings,similarities)]
        return itemAverages[item] + sum(weightedRatings) / sum(similarities)
    else:
        return ratingMean
```

```
In [50]: simPredictions1 = [predictRating1(d['customer_id'], d['product_id']) for d in da
simPredictions2 = [predictRating2(d['customer_id'], d['product_id']) for d in da
simPredictions3 = [predictRating3(d['customer_id'], d['product_id']) for d in da
```

```
In [51]: MSE(alwaysPredictMean, labels)
```

```
Out[51]: 1.4796142779564334
```

```
In [52]: MSE(simPredictions1, labels)
```

```
Out[52]: 1.6146130004291603
```

```
In [53]: MSE(simPredictions2, labels)
```

```
Out[53]: 1.4540822838636853
```

```
In [54]: MSE(simPredictions3, labels)
```

```
Out[54]: 1.44672577948388
```

```
In [ ]:
```