

```
In [1]: import array
import gzip
import matplotlib.pyplot as plt
import random
import tensorflow as tf
from collections import defaultdict
from scipy.spatial import distance
from sklearn.manifold import TSNE
```

Data is available at <http://cseweb.ucsd.edu/~jmcauley/pml/data/>. Download and save to your own directory

```
In [2]: dataDir = "/home/jmcauley/pml_data/"
```

Visual compatibility model

This code reads image data in a specific binary format, described here:

<http://jmcauley.ucsd.edu/data/amazon/links.html>

```
In [3]: def readImageFeatures(path):
        f = open(path, 'rb')
        while True:
            asin = f.read(10).decode('utf-8')
            if len(asin) < 10: break
            a = array.array('f')
            a.fromfile(f, 4096)
            yield str(asin), a.tolist()
```

```
In [4]: def parse(path):
        g = gzip.open(path, 'r')
        for l in g:
            yield eval(l)
```

```
In [5]: X = []
asinPos = {}

for asin, f in readImageFeatures(dataDir + 'image_features_Baby.b'):
    asinPos[asin] = len(X)
    X.append(tf.constant(f, shape=[1, len(f)]))
```

Extract metadata describing ground-truth compatibility relationships among items

```
In [6]: compat = []
asinList = list(asinPos.keys())
for l in parse(dataDir + 'meta_Baby.json.gz'):
    a1 = l['asin']
    if 'related' in l and 'also_bought' in l['related']:
        for a2 in l['related']['also_bought']:
            if a1 in asinPos and a2 in asinPos:
```

```
compat.append((asinPos[a1],asinPos[a2],1))
compat.append((random.randint(0, len(X)-1),random.randint(0, len
```

Number of compatible pairs

```
In [7]: len(compat)
```

```
Out[7]: 1698622
```

```
In [8]: featDim = X[0].shape[1] # Image feature dimensionality
styleDim = 5 # Dimensionality of compressed (projected) representations
```

```
In [9]: featDim
```

```
Out[9]: 4096
```

Define the compatibility model

```
In [10]: optimizer = tf.keras.optimizers.Adam(0.00001)
```

```
In [11]: class CompatibilityModel(tf.keras.Model):
def __init__(self, featDim, styleDim):
    super(CompatibilityModel, self).__init__()
    self.E1 = tf.Variable(tf.random.normal([featDim,styleDim],stddev=0.001))
    self.E2 = tf.Variable(tf.random.normal([featDim,styleDim],stddev=0.001))
    self.c = tf.Variable(0.0)

def predict(self, x1, x2):
    s1 = tf.matmul(x1, self.E1)
    s2 = tf.matmul(x2, self.E2)
    return tf.math.sigmoid(self.c - tf.reduce_sum(tf.math.squared_difference

def call(self, x1, x2, y):
    return -tf.math.log(self.predict(x1,x2)*(2*y - 1) - y + 1)

model = CompatibilityModel(featDim, styleDim)
```

```
In [12]: def trainingStep(compat):
with tf.GradientTape() as tape:
    (i1,i2,y) = random.choice(compat)
    x1,x2 = X[i1],X[i2]
    objective = model(x1,x2,y)
    gradients = tape.gradient(objective, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
    return objective.numpy()
```

```
In [13]: for i in range(50000):
obj = trainingStep(compat)
if (i % 5000 == 4999): print("iteration " + str(i+1) + ", objective = " + st

iteration 5000, objective = 0.51231325
```

```

iteration 10000, objective = 0.6379045
iteration 15000, objective = 0.7741432
iteration 20000, objective = 0.7855767
iteration 25000, objective = 0.46546566
iteration 30000, objective = 0.4155722
iteration 35000, objective = 0.7017366
iteration 40000, objective = 0.83796984
iteration 45000, objective = 0.66832584
iteration 50000, objective = 0.41149664

```

Exercises

9.1

For these exercises we use musical instrument data; we do so because (a) it has fine-grained subcategories (e.g. "accessories", "guitars", etc.) which can be used for these exercises; and (b) because it is small. These exercises might ideally be run with a large category of (e.g.) clothing images, though such datasets are larger and more difficult to work with.

First collect the subcategories associated with each item (for use in Exercise 9.3)

```

In [14]: categories = dict()
itemsPerCategory = defaultdict(set)
for l in parse(dataDir + 'meta_Musical_Instruments.json.gz'):
    cats = l['categories'][0]
    if len(cats) < 2:
        continue
    cat = cats[1] # Extract the "second level" (or sub-) category, for products
    categories[l['asin']] = cat
    itemsPerCategory[cat].add(l['asin'])

```

Read image data

```

In [15]: X = []
asinPos = {}
posPerCategory = defaultdict(set)

for asin, f in readImageFeatures(dataDir + 'image_features_Musical_Instruments.b
    if not asin in categories: # Skip items for which we don't have a category
        continue
    asinPos[asin] = len(X)
    posPerCategory[categories[asin]].add(asinPos[asin])
    X.append(tf.constant(f, shape=[1, len(f)]))

```

Extract compatibility relationships. Build our collection of "difficult" negatives consisting of items from the same category.

```

In [16]: compat = []
asinList = list(asinPos.keys())
for l in parse(dataDir + 'meta_Musical_Instruments.json.gz'):
    a1 = l['asin']
    if not a1 in categories:
        continue
    cat = categories[a1]

```

```

if 'related' in l and 'also_bought' in l['related']:
    for a2 in l['related']['also_bought']:
        if not a2 in categories or categories[a2] != cat:
            continue # Only consider positive relations of the same category
        if a1 in asinPos and a2 in asinPos:
            compat.append((asinPos[a1], asinPos[a2], 1))
            negSameCat = random.sample(posPerCategory[cat], 1)[0]
            compat.append((asinPos[a1], negSameCat, 0))

```

In [17]: `len(compat)`

Out[17]: 809528

In [18]: `featDim = X[0].shape[1] # Image feature dimensionality`
`styleDim = 5 # Dimensionality of compressed (projected) representations`

In [19]: `optimizer = tf.keras.optimizers.Adam(0.00001)`

In [20]: `class CompatibilityModel(tf.keras.Model):`
 `def __init__(self, featDim, styleDim):`
 `super(CompatibilityModel, self).__init__()`
 `self.E1 = tf.Variable(tf.random.normal([featDim, styleDim], stddev=0.001))`
 `self.E2 = tf.Variable(tf.random.normal([featDim, styleDim], stddev=0.001))`
 `self.c = tf.Variable(0.0)`

 `def predict(self, x1, x2):`
 `s1 = tf.matmul(x1, self.E1)`
 `s2 = tf.matmul(x2, self.E2)`
 `return tf.math.sigmoid(self.c - tf.reduce_sum(tf.math.squared_difference`

 `def call(self, x1, x2, y):`
 `return -tf.math.log(self.predict(x1, x2)*(2*y - 1) - y + 1)`

9.2 / 9.3

Modify the model to compute similarity based on the inner product rather than Euclidean distance

In [21]: `class CompatibilityModelInner(tf.keras.Model):`
 `def __init__(self, featDim, styleDim):`
 `super(CompatibilityModel, self).__init__()`
 `self.E1 = tf.Variable(tf.random.normal([featDim, styleDim], stddev=0.001))`
 `self.E2 = tf.Variable(tf.random.normal([featDim, styleDim], stddev=0.001))`
 `self.c = tf.Variable(0.0)`

 `def predict(self, x1, x2):`
 `s1 = tf.matmul(x1, self.E1)`
 `s2 = tf.matmul(x2, self.E2)`
 `return tf.math.sigmoid(self.c + tf.tensordot(s1, s2, 1))`

 `def call(self, x1, x2, y):`
 `return -tf.math.log(self.predict(x1, x2)*(2*y - 1) - y + 1)`

Compare models based on the inner product and Euclidean distance. Both make use of "difficult" negatives (Exercise 9.3)

```
In [22]: model1 = CompatibilityModel(featDim, styleDim)
        model2 = CompatibilityModel(featDim, styleDim)
```

```
In [23]: def trainingStep(model, compat):
        with tf.GradientTape() as tape:
            (i1,i2,y) = random.choice(compat)
            x1,x2 = X[i1],X[i2]
            objective = model(x1,x2,y)
            gradients = tape.gradient(objective, model.trainable_variables)
            optimizer.apply_gradients(zip(gradients, model.trainable_variables))
        return objective.numpy()
```

```
In [24]: random.shuffle(compat)
```

```
In [25]: compatTrain = compat[:700000]
        compatTest = compat[700000:]
```

```
In [26]: for i in range(50000):
        obj = trainingStep(model1, compat)
        if (i % 5000 == 4999): print("iteration " + str(i+1) + ", objective = " + str(obj))

iteration 5000, objective = 0.83570564
iteration 10000, objective = 0.84522265
iteration 15000, objective = 0.6948385
iteration 20000, objective = 0.55048937
iteration 25000, objective = 0.53180134
iteration 30000, objective = 0.7649663
iteration 35000, objective = 0.30317622
iteration 40000, objective = 0.52734387
iteration 45000, objective = 1.1975294
iteration 50000, objective = 0.3939314
```

```
In [27]: for i in range(50000):
        obj = trainingStep(model2, compat)
        if (i % 5000 == 4999): print("iteration " + str(i+1) + ", objective = " + str(obj))

iteration 5000, objective = 0.71674216
iteration 10000, objective = 1.2483406
iteration 15000, objective = 0.6968014
iteration 20000, objective = 0.3971393
iteration 25000, objective = 0.5065777
iteration 30000, objective = 0.75319344
iteration 35000, objective = 0.77804667
iteration 40000, objective = 0.4195373
iteration 45000, objective = 0.39572006
iteration 50000, objective = 0.7114118
```

Compute accuracy (what fraction of positive relationships were predicted as positive)

```
In [28]: acc = 0

for (i1,i2,y) in compatTest:
    x1,x2 = X[i1],X[i2]
    p = model1(x1,x2,y)
    if (p.numpy() > 0.5) == (y == 1):
        acc += 1

acc / len(compatTest)
```

Out[28]: 0.8074738879555913

```
In [29]: acc = 0

for (i1,i2,y) in compatTest:
    x1,x2 = X[i1],X[i2]
    p = model2(x1,x2,y)
    if (p.numpy() > 0.5) == (y == 1):
        acc += 1

acc / len(compatTest)
```

Out[29]: 0.8218994229785991

9.4

t-SNE embedding

```
In [30]: Xembed = []

for asin in asinList:
    i = asinPos[asin]
    x = X[i]
    embedded = list(tf.matmul(x, model1.E1).numpy()[0])
    Xembed.append(embedded)
```




```
In [31]: Xembed2 = TSNE(n_components=2).fit_transform(Xembed)
```

```
In [32]: scatterPlotsX = defaultdict(list)
scatterPlotsY = defaultdict(list)
for xy, asin in zip(Xembed2, asinList):
    if asin in categories:
        cat = categories[asin]
        try:
            scatterPlotsX[cat].append(xy[0])
            scatterPlotsY[cat].append(xy[1])
        except Exception as e:
            pass
```

Scatterplots by subcategory aren't particularly interesting in this case. Try e.g. price or brand for more compelling examples.

```
In [33]: for cat in ['Instrument Accessories', 'Stringed Instruments', 'Guitars']:
```

```
plt.scatter(scatterPlotsX[cat],
            scatterPlotsY[cat], color='lightgrey', lw = 0, label = cat)
plt.legend(loc='lower left')
plt.xticks([])
plt.yticks([])
plt.xlabel("first embedded dimension ")
plt.ylabel("second embedded dimension")
plt.title("\emph{TSNE}-based item embeddings")
plt.show()
```



In []: